

# Lab4: spark MLlib

## Task 1-2: Load Data & Show Schema

- **Relation to Course:** Basic Spark operations to ingest data (prerequisite for MLlib workflows).
- **Code:**

```
df = spark.read.csv("tp4_data/*", header=True, inferSchema=True)
df.printSchema()
```

## Task 3: Fill Missing Values

- **Relation:** Data cleaning (similar to preprocessing in `RFormula`).
- **Code:**

```
df = df.fillna(0)
```

## Task 4: Add "day\_of\_week" Column

- **Relation:** Feature engineering using Spark SQL functions.
- **Code:**

```
from pyspark.sql.functions import date_format
df = df.withColumn("day_of_week", date_format("InvoiceDate", "EEEE"))
```

## Task 5: Temporal Split

- **Relation:** Data splitting for training/testing (non-random split, unlike the course example).
- **Code:**

```
train = df.filter(df.InvoiceDate < "2010-12-13")
test = df.filter(df.InvoiceDate >= "2010-12-13")
```

## Task 6-7: StringIndexer & Ordinal Issue

- **Relation:** Encoding categorical variables manually (vs. `RFormula` automation).
- **Code:**

```
from pyspark.ml.feature import StringIndexer
indexer = StringIndexer(inputCol="day_of_week", outputCol="day_of_week_encoded")
```

- **Ordinal Issue:** `StringIndexer` assigns arbitrary numerical labels (e.g., Monday=0, Tuesday=1). To avoid implying order, use `OneHotEncoder`:

```
from pyspark.ml.feature import OneHotEncoder
encoder = OneHotEncoder(inputCol="day_of_week_encoded", outputCol="day_of_week_vec")
```

## Task 8: VectorAssembler

- **Relation:** Combining features into a vector (like `RFormula`'s output).
- **Code:**

```
from pyspark.ml.feature import VectorAssembler
assembler = VectorAssembler(
    inputCols=["UnitPrice", "Quantity", "day_of_week_encoded"],
    outputCol="features"
)
```

## Task 9: Pipeline

- **Relation:** Chains stages (transformers + estimator) like the course example.
- **Code:**

```
from pyspark.ml import Pipeline
pipeline = Pipeline(stages=[indexer, assembler, kmeans])
```

- **The purpose:**

The pipeline consists of **multiple stages** (e.g., `StringIndexer`, `OneHotEncoder`, `VectorAssembler`). When you **fit** the pipeline on the training set, it learns transformations (e.g., how to encode categorical variables). Then, you can **reuse** the same pipeline to transform new data (like the test set) **without redefining everything manually**.

## Task 10: Unique Values in StringIndexer

- **Relation:** `StringIndexer` must `fit()` to learn categories (like `RFormula`'s `fit()` step).
- **Solution:** Call `fit()` on the training data to ensure all categories are seen:

```
indexer_model = indexer.fit(train)
```

## Task 11-14: Transform & Train KMeans

- **Relation:** Train a model (estimator) using transformed data.
- **Code:**

```
model = pipeline.fit(train) # Includes KMeans training
test_predictions = model.transform(test)
```

## Task 15: Silhouette Coefficient

- **Relation:** Evaluator for clustering (like `BinaryClassificationEvaluator` for classification).
- **Code:**

```
from pyspark.ml.evaluation import ClusteringEvaluator
evaluator = ClusteringEvaluator()
silhouette = evaluator.evaluate(test_predictions)
print(f"Silhouette: {silhouette}")
```